

# AI/ML for Climate Workshop

International Livestock Research Institute (ILRI)

hide: - toc

## ML-based Subseasonal & Seasonal Prediction

In this section we explore the prepared predictors and target data:

- Are there missing values?
- What are the distributions and ranges?
- Do we see relationships between drivers (predictors) and rainfall anomalies (target)?

## Interactive Learning



**Click the Binder button above to launch an interactive Jupyter notebook for NumPy and Pandas climate data analysis!**

```
# EDA in Notebook
```

```
import os
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
from scipy.stats import linregress

PROCESSED_DIR = "data/processed"
REPORT_FIGS   = "reports/figures"
REPORT_TBLS   = "reports/tables"
os.makedirs(PROCESSED_DIR, exist_ok=True)
os.makedirs(REPORT_FIGS, exist_ok=True)
```

```
os.makedirs(REPORT_TBLS, exist_ok=True)

# Set working directory

os.chdir("C:\\Users\\yonas\\Documents\\ICPAC\\ea_seasonal_pred\\seasonal-ml-pred")
```

```
# Import Uganda Data

data_uganda_ond = pd.read_csv("data/processed/uganda_OND_target_features_1981_2024.csv")
data_uganda_ond.set_index("season_year", inplace=True)

data_uganda_ond.head()
```

### Output:

season_year	OND_total_mm	OND_clim1991_2020_mm	OND_anom_mm	OND_anom_std	\
1981-01-01	251.31606	328.4043	-77.088240	-1.053309	
1982-01-01	373.33360	328.4043	44.929290	0.613899	
1983-01-01	269.18472	328.4043	-59.219574	-0.809157	
1984-01-01	294.98053	328.4043	-33.423767	-0.456691	
1985-01-01	262.54230	328.4043	-65.862000	-0.899917	

season_year	n12_JAS	n12_AS	n12_Sep	n3_JAS	n3_AS	n3_Sep	...	\
1981-01-01	-0.936667	-0.960	-0.79	-0.520000	-0.480	-0.23	...	
1982-01-01	0.860000	1.055	1.31	1.083333	1.420	1.89	...	
1983-01-01	2.700000	2.150	1.37	0.526667	0.390	0.13	...	
1984-01-01	-0.420000	-0.315	-0.07	-0.510000	-0.375	-0.34	...	
1985-01-01	-1.206667	-1.110	-0.99	-0.946667	-0.870	-0.86	...	

season_year	pacwarmpool_Sep	censo_JAS	censo_AS	censo_Sep	dmi_JAS	\
1981-01-01	-0.186	-0.360000	-0.235	-0.18	-0.648333	
1982-01-01	-0.495	1.786667	1.905	1.95	0.321000	
1983-01-01	-0.082	-0.070000	-0.285	-0.64	0.267000	
1984-01-01	-0.288	-0.250000	-0.265	-0.21	-0.490667	
1985-01-01	-0.381	-0.390000	-0.460	-0.28	-0.366667	

season_year	dmi_AS	dmi_Sep	PERSIST_MJJ_anom_mm	PERSIST_JJA_anom_mm	\
1981-01-01	-0.6925	-0.757	-2.042121	1.513391	
1982-01-01	0.3490	0.442	-8.787225	-1.924204	
1983-01-01	0.1380	-0.069	-0.441927	2.636551	
1984-01-01	-0.5530	-0.608	-10.111957	-6.647986	
1985-01-01	-0.3485	-0.238	-4.830346	0.538307	

season_year	PERSIST_JAS_anom_mm
-------------	---------------------

```

1981-01-01      3.086952
1982-01-01     -5.021202
1983-01-01      2.906827
1984-01-01     -6.963177
1985-01-01      1.014789

```

```
[5 rows x 34 columns]
```

```

def ensure_year_index(df: pd.DataFrame, name="season_year") -> pd.DataFrame:
    df = df.copy()
    if isinstance(df.index, pd.DatetimeIndex):
        df.index = df.index.year
    else:
        try:
            df.index = df.index.astype(int)
        except Exception:
            df.index = pd.to_datetime(df.index, errors="coerce").year
    df.index = pd.Index(df.index, name=name)
    return df

```

```

data_uganda_ond = ensure_year_index(data_uganda_ond)
data_uganda_ond.head()

```

### Output:

```

      OND_total_mm  OND_clim1991_2020_mm  OND_anom_mm  OND_anom_std  \
season_year
1981           251.31606             328.4043   -77.088240   -1.053309
1982           373.33360             328.4043    44.929290    0.613899
1983           269.18472             328.4043   -59.219574   -0.809157
1984           294.98053             328.4043   -33.423767   -0.456691
1985           262.54230             328.4043   -65.862000   -0.899917

      n12_JAS  n12_AS  n12_Sep   n3_JAS  n3_AS  n3_Sep  ...  \
season_year
1981    -0.936667  -0.960    -0.79  -0.520000  -0.480   -0.23  ...
1982     0.860000   1.055     1.31   1.083333   1.420    1.89  ...
1983     2.700000   2.150     1.37   0.526667   0.390    0.13  ...
1984    -0.420000  -0.315    -0.07  -0.510000  -0.375   -0.34  ...
1985    -1.206667  -1.110    -0.99  -0.946667  -0.870   -0.86  ...

      pacwarmpool_Sep  censo_JAS  censo_AS  censo_Sep  dmi_JAS  \
season_year
1981          -0.186  -0.360000   -0.235    -0.18  -0.648333
1982          -0.495   1.786667    1.905     1.95   0.321000
1983          -0.082  -0.070000   -0.285    -0.64   0.267000
1984          -0.288  -0.250000   -0.265    -0.21  -0.490667
1985          -0.381  -0.390000   -0.460    -0.28  -0.366667

      dmi_AS  dmi_Sep  PERSIST_MJJ_anom_mm  PERSIST_JJA_anom_mm  \

```

```
season_year
1981      -0.6925   -0.757          -2.042121          1.513391
1982       0.3490    0.442          -8.787225         -1.924204
1983       0.1380   -0.069          -0.441927          2.636551
1984      -0.5530   -0.608         -10.111957         -6.647986
1985      -0.3485   -0.238          -4.830346          0.538307
```

PERSIST\_JAS\_anom\_mm

```
season_year
1981              3.086952
1982             -5.021202
1983              2.906827
1984             -6.963177
1985              1.014789
```

[5 rows x 34 columns]

```
def year_array(idx) -> np.ndarray:
    """Return an array of year ints from either DatetimeIndex or int-like index."""
    if isinstance(idx, pd.DatetimeIndex):
        return idx.year.to_numpy(dtype=int)
    # try casting to int directly
    try:
        return pd.Index(idx).astype(int).to_numpy()
    except Exception:
        # last resort: parse as datetime and take year
        return pd.to_datetime(idx, errors="coerce").year.to_numpy(dtype=int)
```

```
REPORT_FIGS = "reports/figures"
REPORT_TBLS = "reports/tables"
os.makedirs(REPORT_FIGS, exist_ok=True)
os.makedirs(REPORT_TBLS, exist_ok=True)

def summarize_target(table, season: str, anomaly_col: str, total_col: str, clim_col: str):
    y = table[anomaly_col].copy()
    tot = table[total_col].copy()
    clim = table[clim_col].copy()

    # --- Summary stats
    desc = y.describe(percentiles=[.05, .10, .25, .50, .75, .90, .95]).to_frame("value")
    desc.to_csv(os.path.join(REPORT_TBLS, f"{season}_target_anomaly_summary.csv"))
    print(f"=== {season} target summary (anomaly mm) ===")
    display(desc)

    # --- Time series
    fig, ax = plt.subplots(figsize=(10,4))
    ax.plot(y.index, y.values, label="Anomaly (mm)")
    ax.axhline(0, linestyle="--", linewidth=1)
    ax.set_title(f"{season}: Seasonal rainfall anomaly")
    ax.set_xlabel("Year"); ax.set_ylabel("mm"); ax.legend()
    fig.tight_layout(); fig.savefig(os.path.join(REPORT_FIGS, f"{season}_anomaly_ts.png"))
```

```

plt.show()

# --- Totals vs climatology
fig, ax = plt.subplots(figsize=(10,4))
ax.plot(tot.index, tot.values, label="Total (mm)")
ax.plot(clim.index, clim.values, label="Climatology 1991-2020 (mm)")
ax.set_title(f"{season}: Totals vs climatology")
ax.set_xlabel("Year"); ax.set_ylabel("mm"); ax.legend()
fig.tight_layout(); fig.savefig(os.path.join(REPORT_FIGS, f"{season}_totals_vs_clim
plt.show()

# --- Decadal boxplot (robust to datetime index)
df = y.dropna().to_frame("anom")
yrs = year_array(df.index) # <-- convert index to years here
df["decade"] = (yrs // 10) * 10
decades = np.unique(df["decade"].astype(int))
data = [df.loc[df["decade"] == d, "anom"].values for d in decades]

fig, ax = plt.subplots(figsize=(8,4))
ax.boxplot(data, tick_labels=[str(int(d)) for d in decades], showmeans=True)
ax.set_title(f"{season}: Anomaly by decade")
ax.set_xlabel("Decade")
ax.set_ylabel("mm")
fig.tight_layout()
fig.savefig(os.path.join(REPORT_FIGS, f"{season}_anomaly_by_decade.png"), dpi=150)
plt.show()

# --- Linear trend (robust)
y_no_na = y.dropna()
yr_num = year_array(y_no_na.index).astype(float)
yy = y_no_na.values
if len(yy) > 3:
    slope, intercept, r, p, se = linregress(yr_num, yy)
    print(f"Trend {season}: slope={slope:.2f} mm/yr (p={p:.3f}, r={r:.2f})")
else:
    print("Not enough data for trend.")

```

```

# Example with your Uganda OND table (must include these columns)
# data_uganda_ond.index can be datetime or ints; both are fine now.
summarize_target(
    data_uganda_ond,
    season="OND",
    anomaly_col="OND_anom_mm",
    total_col="OND_total_mm",
    clim_col="OND_clim1991_2020_mm",
)

```

**Output:**

```

=== OND target summary (anomaly mm) ===

      value
count  44.000000
mean   -4.571462
std     66.803853
min    -121.979034
5%     -87.610841
10%    -73.110354
25%    -49.469459
50%    -17.898682
75%     33.797585
90%     82.521012
95%    113.560480
max     190.410400
<Figure size 1000x400 with 1 Axes>
<Figure size 1000x400 with 1 Axes>
<Figure size 800x400 with 1 Axes>
Trend OND: slope=1.46 mm/yr (p=0.064, r=0.28)

```

```

from scipy.stats import pearsonr, spearmanr

```

```

def corr_with_targets(df: pd.DataFrame, target_cols: list[str], feature_cols: list[str])
    """
    Pairwise correlations between each feature and each target.
    Returns (R, P, N) as DataFrames (rows=features, cols=targets).
    """
    R = pd.DataFrame(index=feature_cols, columns=target_cols, dtype=float)
    P = pd.DataFrame(index=feature_cols, columns=target_cols, dtype=float)
    N = pd.DataFrame(index=feature_cols, columns=target_cols, dtype=int)

    stat = pearsonr if method.lower().startswith("pear") else spearmanr

    for t in target_cols:
        for f in feature_cols:
            sub = df[[t, f]].dropna()
            n = len(sub)
            N.loc[f, t] = n
            if n >= 3:
                r, p = stat(sub[t].values, sub[f].values)
                R.loc[f, t] = r
                P.loc[f, t] = p
            else:
                R.loc[f, t] = np.nan
                P.loc[f, t] = np.nan
    return R, P, N

def heatmap_feature_vs_target(R: pd.DataFrame, P: pd.DataFrame, title: str, outfile_png: str)
    """
    Plot heatmap of correlations: rows=features, cols=targets.
    sort_by: "abs" or "value" to sort rows by correlation with target_for_sort.

```

```

"""
rmat = R.copy()

# sort features (rows)
if sort_by is not None and target_for_sort in rmat.columns:
    key = rmat[target_for_sort].abs() if sort_by == "abs" else rmat[target_for_sort]
    order = key.sort_values(ascending=False).index.tolist()
    rmat = rmat.loc[order]

# limit to topn rows for readability
if topn is not None and rmat.shape[0] > topn:
    rmat = rmat.iloc[:topn, :]

# plot
nrows, ncols = rmat.shape
fig_h = max(4, 0.35 * nrows + 1) # dynamic height
fig_w = max(4, 1.2 * ncols + 2)
fig, ax = plt.subplots(figsize=(fig_w, fig_h))
im = ax.imshow(rmat.values, vmin=-1, vmax=1) # default colormap
cbar = fig.colorbar(im, ax=ax, fraction=0.046, pad=0.04)
cbar.set_label("Correlation (r)")

ax.set_yticks(range(nrows)); ax.set_yticklabels(rmat.index, fontsize=8)
ax.set_xticks(range(ncols)); ax.set_xticklabels(rmat.columns, rotation=0, fontsize=8)
ax.set_title(title)
ax.set_xlabel("Target"); ax.set_ylabel("Feature")

# optional text annotations with r and significance stars
if annotate:
    # significance: * (p<=0.05), ** (p<=0.01), *** (p<=0.001)
    for i in range(nrows):
        for j in range(ncols):
            r = rmat.iat[i, j]
            p = P.loc[rmat.index[i], rmat.columns[j]]
            if np.isfinite(r):
                stars = "****" if p is not None and p <= 0.001 else ("***" if p is not None and p <= 0.01 else "**" if p is not None and p <= 0.05 else "")
                ax.text(j, i, f"{r: .2f}{stars}", ha="center", va="center", fontsize=8)

fig.tight_layout()
fig.savefig(outfile_png, dpi=150)
plt.show()

```

```

def heatmap_feature_vs_target(
    R: pd.DataFrame,
    P: pd.DataFrame,
    title: str,
    outfile_png: str,
    topn: int = 40,
    sort_by: str | None = None,
    target_for_sort: str | None = None,
    annotate: bool = True,
    figsize: tuple[float, float] | None = None,

```

```

ytick_fontsize: int = 10,
xtick_fontsize: int = 11,
left_margin: float = 0.35, # more space for long feature names
):
    """
    Plot heatmap of correlations: rows=features, cols=targets.
    figsize: (width, height) inches. If None, a sensible wide default is used.
    """
    rmat = R.copy()

    # sort rows by correlation with the chosen target
    if sort_by is not None and target_for_sort in rmat.columns:
        key = rmat[target_for_sort].abs() if sort_by == "abs" else rmat[target_for_sort]
        rmat = rmat.loc[key.sort_values(ascending=False).index]

    # limit to topn rows
    if topn is not None and rmat.shape[0] > topn:
        rmat = rmat.iloc[:topn, :]

    nrows, ncols = rmat.shape

    # --- wide default if not provided
    if figsize is None:
        fig_h = max(5, 0.32 * nrows + 2.0)
        fig_w = max(12, 3.0 * ncols + 8.0) # force a wide figure even with 1 target c
        figsize = (fig_w, fig_h)

    fig, ax = plt.subplots(figsize=figsize)
    im = ax.imshow(rmat.values, vmin=-1, vmax=1)
    cbar = fig.colorbar(im, ax=ax, fraction=0.046, pad=0.04)
    cbar.set_label("Correlation (r)")

    ax.set_yticks(range(nrows)); ax.set_yticklabels(rmat.index, fontsize=ytick_fontsize)
    ax.set_xticks(range(ncols)); ax.set_xticklabels(rmat.columns, fontsize=xtick_fontsize)
    ax.set_title(title); ax.set_xlabel("Target"); ax.set_ylabel("Feature")

    # annotate with r and significance stars
    if annotate:
        for i in range(nrows):
            for j in range(ncols):
                r = rmat.iat[i, j]
                p = P.loc[rmat.index[i], rmat.columns[j]]
                if np.isfinite(r):
                    stars = "" if p is not None and p <= 0.001 else ("**" if p is not None else "")
                    ax.text(j, i, f"{r: .2f}{stars}", ha="center", va="center",
                            fontsize=8, color="white" if abs(r) > 0.5 else "black")

    # give labels breathing room on the left
    fig.subplots_adjust(left=left_margin, right=0.98, top=0.92, bottom=0.08)
    fig.savefig(outfile_png, dpi=150)
    plt.show()

```



```
heatmap_feature_vs_target(
    R_ond, P_ond,
    title="OND: Feature ↔ Target correlation (Pearson)",
    outfile_png=os.path.join(REPORT_FIGS, "OND_feature_target_heatmap_wide.png"),
    topn=40,
    sort_by="abs",
    target_for_sort="OND_anom_mm",
    annotate=True,
    figsize=(16, 18),          # ← make it wider and taller
    ytick_fontsize=10,
    left_margin=0.42          # ← more space for long feature names
)
```

### Output:

<Figure size 1600x1800 with 2 Axes>

```
heatmap_feature_vs_target(
    R_mam, P_mam,
    title="MAM: Feature ↔ Target correlation (Pearson)",
    outfile_png=os.path.join(REPORT_FIGS, "MAM_feature_target_heatmap_wide.png"),
    topn=40,
    sort_by="abs",
    target_for_sort="MAM_anom_mm",
    annotate=True,
    figsize=(14, 18),
    ytick_fontsize=10,
    left_margin=0.42
)
```

```
def load_model_table(parq: str, csv: str):
    p = os.path.join(PROCESSED_DIR, parq)
    c = os.path.join(PROCESSED_DIR, csv)
    if os.path.exists(p):
        df = pd.read_parquet(p)
    elif os.path.exists(c):
        df = pd.read_csv(c, index_col=0)
    else:
        raise FileNotFoundError(f"Missing {p} or {c}")
    # Ensure year index (int)
    if isinstance(df.index, pd.DatetimeIndex):
        df.index = df.index.year
    else:
        try:
            df.index = df.index.astype(int)
        except Exception:
            df.index = pd.to_datetime(df.index, errors="coerce").year.astype(int)
    df.index.name = "season_year"
    return df.sort_index()
```

```

OND_targets = ["OND_anom_mm", "OND_total_mm", "OND_anom_std"]
MAM_targets = ["MAM_anom_mm", "MAM_total_mm", "MAM_anom_std"]

OND_features = [c for c in data_uganda_ond.columns if c not in ["OND_total_mm", "OND_cli
#MAM_features = [c for c in data_uganda_mam.columns if c not in ["MAM_total_mm", "MAM_cl

def zscore(s: pd.Series) -> pd.Series:
    m, sd = s.mean(), s.std(ddof=1)
    return (s - m) / sd if sd and not np.isclose(sd, 0.0) else s * np.nan

def rolling_if(s: pd.Series, win: int) -> pd.Series:
    return s if win in (0, 1) else s.rolling(win, min_periods=max(2, win//2)).mean()

```

```

import ipywidgets as widgets
from IPython.display import display, clear_output

```

```

season_dd = widgets Dropdown(options=["OND", "MAM"], value="OND", description="Season:")
scale_dd = widgets Dropdown(options=["z-score overlay", "raw (twin axis)"], value="z-sc
smooth_sl = widgets.IntSlider(value=0, min=0, max=7, step=1, description="Roll mean (yr

# These will be (re)built when season changes
target_dd = widgets Dropdown(description="Target:")
feat_ms = widgets.SelectMultiple(description="Features:", rows=12, layout=widgets.Lay

year_range = widgets.IntRangeSlider(description="Year range:", min=1980, max=2025, step
                                continuous_update=False, layout=widgets.Layout(widt

out = widgets.Output()

```

```

def setup_for_season(season: str):
    if season == "OND":
        df = OND; t_opts = OND_targets; f_opts = OND_features
    else:
        df = MAM; t_opts = MAM_targets; f_opts = MAM_features

    target_dd.options = [c for c in t_opts if c in df.columns]
    target_dd.value = target_dd.options[0]

    # Suggest a few default features if present
    defaults = {
        "OND": ["N34_JAS", "DMI_JAS", "PERSIST_JAS_anom_mm"],
        "MAM": ["N34_NDJ", "DMI_NDJ", "PERSIST_DJF_anom_mm"]
    }
    feat_opts = [c for c in f_opts if c in df.columns]
    feat_ms.options = feat_opts

```

```

    feat_ms.value = tuple([c for c in defaults[season] if c in feat_opts])

    # Year slider bounds
    year_min, year_max = int(df.index.min()), int(df.index.max())
    year_range.min, year_range.max = year_min, year_max
    year_range.value = (max(1981, year_min), year_max)

    setup_for_season(season_dd.value)

    def on_season_change(change):
        if change["name"] == "value":
            setup_for_season(change["new"])
    season_dd.observe(on_season_change)

    def plot_interactive(*args):
        out.clear_output()
        with out:
            season = season_dd.value
            df = OND if season == "OND" else MAM
            ycol = target_dd.value
            fcols = list(feat_ms.value)

            if not fcols:
                print("Select at least one feature.")
                return

            y0, y1 = year_range.value
            df = df.loc[(df.index >= y0) & (df.index <= y1)].copy()

            # Build series (align and drop NaNs per pair)
            y = df[ycol].astype(float)
            features = {c: df[c].astype(float) for c in fcols}

            # Optional smoothing
            win = int(smooth_sl.value)
            y = rolling_if(y, win)
            features = {k: rolling_if(v, win) for k, v in features.items()}

            # Plot
            fig, ax = plt.subplots(figsize=(12, 5))
            ax.set_title(f"{season}: {ycol} with selected features")

            if scale_dd.value == "z-score overlay":
                y_plot = zscore(y)
                ax.plot(y_plot.index, y_plot.values, linewidth=2.2, label=ycol)
                for k, s in features.items():
                    ax.plot(s.index, zscore(s).values, label=k, alpha=0.9)
                ax.set_ylabel("z-score")
            else:
                # raw target on left, features on right (twin)
                ax.plot(y.index, y.values, linewidth=2.2, label=ycol, color="k")
                ax.set_ylabel(ycol)
                ax2 = ax.twinx()
                for k, s in features.items():

```

```

        ax2.plot(s.index, s.values, label=k, alpha=0.9)
    ax2.set_ylabel("features (raw units)")
    # merge legends
    lines1, labels1 = ax.get_legend_handles_labels()
    lines2, labels2 = ax2.get_legend_handles_labels()
    ax2.legend(lines1+lines2, labels1+labels2, loc="upper left")
    if scale_dd.value == "z-score overlay":
        ax.legend(loc="upper left")
    ax.set_xlabel("Year")
    ax.grid(True, alpha=0.3)
    plt.show()

# Correlation table (pairwise complete)
rows = []
for k, s in features.items():
    sub = pd.concat([y, s], axis=1).dropna()
    if len(sub) >= 3:
        r = sub.iloc[:,0].corr(sub.iloc[:,1])
        rows.append({"feature": k, "r_with_target": r, "N": len(sub)})
    else:
        rows.append({"feature": k, "r_with_target": np.nan, "N": len(sub)})
corr_df = pd.DataFrame(rows).sort_values("r_with_target", key=np.abs, ascending=False)
display(corr_df.reset_index(drop=True))

```

```

ui = widgets.VBox([
    widgets.HBox([season_dd, scale_dd, smooth_sl]),
    widgets.HBox([target_dd]),
    widgets.HBox([feat_ms]),
    widgets.HBox([year_range]),
])

for w in [season_dd, scale_dd, smooth_sl, target_dd, feat_ms, year_range]:
    w.observe(lambda change: plot_interactive(), names="value")

display(ui, out)
plot_interactive() # initial draw

```