

AI/ML for Climate Workshop

International Livestock Research Institute (ILRI)

hide: - toc

Introduction to Machine Learning for Weather & Climate

1. Learning Goals

By the end of this module, you will be able to:

- Understand what “machine learning” means in an operational climate / meteorological context.
 - Recognize common supervised ML tasks in climate science (regression, classification).
 - Explain the difference between training, validation, and testing.
 - Describe typical inputs (predictors/features) and outputs (predictands/targets) in seasonal and subseasonal prediction.
 - Train a simple ML model in Python using `scikit-learn`.
-

2. Why ML in Climate and Meteorology?

Traditional NWP (Numerical Weather Prediction) and physics-based seasonal models are powerful, but they have limitations:

- **Biases and systematic errors** in forecasts.
- **Resolution gaps** (global models are often too coarse for local decisions).
- **Cost / compute constraints** for running large ensembles in real time.

Machine learning can help with:

- **Bias correction & calibration**
e.g. adjusting a raw seasonal forecast to better match observed rainfall over the Greater Horn of Africa.
- **Downscaling / localization**
e.g. taking coarse $1^\circ \times 1^\circ$ model output and predicting station-scale rainfall probability.
- **Predicting impacts, not just weather**
e.g. “likelihood of below-normal rainfall affecting crop yield in eastern Ethiopia.”

⚠ *ML does **not** replace physics.*

It helps learn statistical relationships between inputs (e.g. SST, winds, geopotential height) and outcomes (e.g. rainfall tercile).

3. Core ML Terms (Climate Context)

Feature / Predictor

A variable we use as input to the model (e.g. Niño3.4 index, soil moisture anomaly, 850 hPa wind, previous month's rainfall).

Target / Label / Predictand

What we're trying to predict (e.g. rainfall category: below/normal/above for OND season in the GHA).

Regression vs Classification

- **Regression:** predict a number
Example: “Forecast total mm of rainfall for next week in Addis Ababa.”
- **Classification:** predict a category
Example: “Is next OND season likely to be ‘below normal’, ‘normal’, or ‘above normal’?”

Training / Validation / Testing

- **Training set:** we fit the model here.
- **Validation set:** we tune hyperparameters here.
- **Test set:** we evaluate final skill here.

In climate, we often split data **by time** (e.g. train on 1993–2015, test on 2016–2024) to simulate “forecasting the future.”

4. Typical ML Workflow for Seasonal/Subseasonal Prediction

1. Data collection

2. Observations / reanalysis (e.g. CHIRPS rainfall, ERA5 winds).
3. Model forecasts / ensemble members (e.g. GCM seasonal forecasts).
4. Climate indices (e.g. IOD, ENSO).

5. Preprocessing

6. Compute anomalies (remove the climatological mean).
7. Spatial averaging (e.g. average rainfall over a GHA subregion).
8. Temporal aggregation (e.g. last 30 days, last 90 days).

9. Feature engineering

10. Build meaningful predictors, like:
 - Regional mean SST anomaly (western Indian Ocean)
 - Zonal wind shear
 - Previous month rainfall deficit
11. Scale / normalize as needed.

12. Model training

13. Fit a regression or classification model.

14. Evaluation

15. Use skill metrics (correlation, RMSE, ACC, Brier score, ROC AUC, etc.).
16. Check reliability and sharpness (is the forecast probabilistic and trustworthy?).

17. Deployment

18. Generate an operational forecast for “next period.”
19. Communicate uncertainty.

5. A Minimal Hands-On Example

Below is a **toy** supervised learning example in `scikit-learn`.

The goal: predict rainfall anomaly (mm/day) from a few climate indices.

In the real workshop, you will replace the dummy `x` and `y` arrays with: - Predictors from reanalysis / model output - Target from observed rainfall over your domain

```
import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score, mean_squared_error

# -----
# Example: X are "features"
#   col 0: ENSO index (e.g. Niño3.4 SST anomaly)
#   col 1: Indian Ocean Dipole index
#   col 2: previous month's regional rainfall anomaly
#
# y is the rainfall anomaly we're trying to predict for next month
# -----

# Fake data for illustration only
X = np.array([
    [1.2, -0.3,  5.1],
    [0.8, -0.1,  4.7],
    [0.1,  0.4, -2.0],
    [-0.5, 0.9, -3.1],
    [-1.0, 1.2, -4.0],
    [0.4, -0.2,  2.2],
    [1.1, -0.4,  4.9],
    [-0.8, 1.1, -3.6],
])

y = np.array([
    6.2,
    5.9,
    -1.8,
    -2.5,
    -3.2,
    2.4,
    5.8,
    -2.9
])

# Split train/test (for demo we'll just cut in half)
X_train, X_test = X[:5], X[5:]
y_train, y_test = y[:5], y[5:]

model = LinearRegression()
model.fit(X_train, y_train)

y_pred = model.predict(X_test)

print("Predictions:", y_pred)
print("Truth:", y_test)
```

```
print("R² score:", r2_score(y_test, y_pred))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred)))
```

Discussion:

- Does the model capture the sign (wet vs dry anomaly)?
- Is skill (R^2 , RMSE) “good enough” to trust it in operations?
- Which predictors seem most important?

We can inspect learned coefficients:

```
for name, coef in zip(["ENSO", "IOD", "PrevRain"], model.coef_):
    print(f"{name}: {coef:.3f}")

print("Intercept:", model.intercept_)
```

This tells you how strongly each climate signal influences the predicted rainfall anomaly in this simple linear model.

6. Classification Example (Tercile Forecasting)

Seasonal forecast centers often issue tercile forecasts: - **Below Normal** - **Normal** - **Above Normal**

Below is a simple example using `LogisticRegression` for classification.

```
import numpy as np
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report

# Let's say classes are:
# 0 = Below Normal
# 1 = Normal
# 2 = Above Normal
#
# X = predictors (same idea: ENSO, IOD, previous rainfall anomaly)
# y = rainfall category next season

X = np.array([
    [1.2, -0.3, 5.1],
    [0.8, -0.1, 4.7],
    [0.1, 0.4, -2.0],
    [-0.5, 0.9, -3.1],
    [-1.0, 1.2, -4.0],
    [0.4, -0.2, 2.2],
    [1.1, -0.4, 4.9],
    [-0.8, 1.1, -3.6],
```

```

])

y = np.array([2, 2, 1, 0, 0, 1, 2, 0])

X_train, X_test = X[:5], X[5:]
y_train, y_test = y[:5], y[5:]

clf = LogisticRegression(multi_class="multinomial", max_iter=500)
clf.fit(X_train, y_train)

y_pred = clf.predict(X_test)

print("Predicted classes:", y_pred)
print("True classes      :", y_test)
print("Accuracy:", accuracy_score(y_test, y_pred))
print("
Full report:
", classification_report(y_test, y_pred))

```

Operational relevance:

- You can map class probabilities (e.g. 70% chance “below normal”) over the Greater Horn of Africa.
- That turns into a forecast product that NMHSs and ICPAC can communicate.

7. Limitations and Good Practice

Limitations - Small sample sizes (seasonal data = maybe only 30–40 years of hindcasts). - Climate is non-stationary (warming trend, changing teleconnections). - Risk of overfitting if you add too many predictors.

Good practice - Use cross-validation that respects time order. - Always compare against a baseline (e.g. climatology, persistence). - Communicate uncertainty (probabilities, confidence).

8. Exercises / Your Turn

1. **Build your own dataset**
2. Choose 2–4 predictors you believe affect OND rainfall in your country.
3. Construct `x` and `y` arrays from historical data.
4. **Train a regression model**
5. Fit `LinearRegression` and report RMSE.

6. Train a classification model

7. Define categories (below/normal/above).

8. Fit `LogisticRegression` .

9. Report accuracy.

10. Interpretability

11. Which predictors were most important?

12. Does that match known physical drivers (ENSO, IOD, etc.)?

13. **Bonus (if time):**

14. Plot forecast vs observed on a time series.

15. Make a simple reliability diagram.

9. Takeaways

- ML in climate is not magic — it's statistics + domain knowledge.
- The value comes from combining:
 - physically meaningful predictors,
 - careful validation,
 - and clear communication of uncertainty.
- You'll build on this in the next module: **"ML Workflow for Weather & Climate"** (`day4/04-ml-workflow.md`),
which focuses on making this operational and reproducible.